

Лабораторная работ 2. Airflow and docker.

Грибанов Данил, 6233

Ход работы

Все файлы (dockerfile, docker-compose и прочее) находятся в папке airflow (так же dags и скрипты для запуска соотв. частей).

Перед запуском самой системы (докера airflow) следует подготовить следующее:

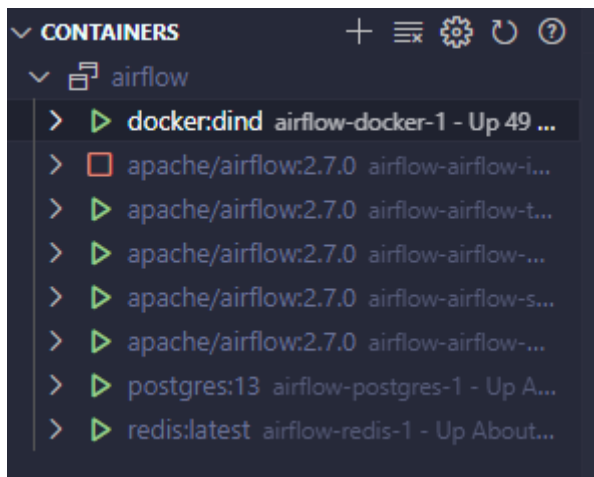
Создадим хранилище с данными для запуска и обучения всех систем:

```
313
314 volumes:
315     postgres-db-volume:
316     airflow-data-volume:
317         driver: local
318         driver_opts:
319             type: none
320             o: bind
321             device: "${AIRFLOW_PROJ_DIR:-.}/data"
```

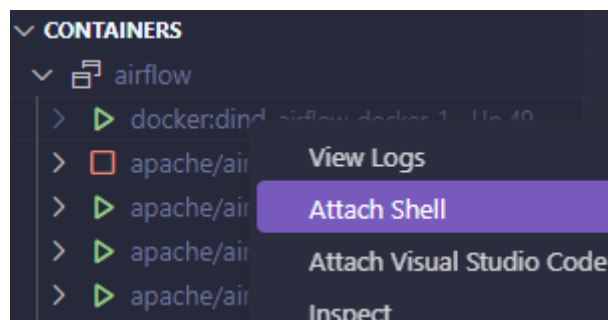
И добавим данное хранилище с папкой докер-файлов нужных в dind:

```
services:
  docker:
    image: docker:dind
    privileged: true
    # Build additional dockers in dind TODO: Is there some way to do init prebuild for dind?
    # command: 'cd /dockerfiles/huggy-face && docker build . -t huggy_face_image'
    environment:
      DOCKER_TLS_CERTDIR: ""
    volumes:
      - airflow-data-volume:/data
      - ${AIRFLOW_PROJ_DIR:-.}/dockerfiles:/dockerfiles
  postgres:
    image: postgres:13
```

Запустим docker-compose.yaml для сборки airflow и нужного:



Дождемся инициализации (пару минут), и зайдем в docker:dind:



Теперь подготовим здесь докер для использования нейронных сетей:

```
cd /dockerfiles/huggy-face && docker build . -t huggy_face_image
```

Дождемся сборки (до 10 минут). Теперь мы готовы к запуску основных DAG разработанных для лабораторных работ. Код DAG к соотв. пунктам задания имеет комментарии и пояснения.

Видео в краткий пересказ

Airflow/dags/airflow_lab2.py

```
1. import os
2. from datetime import datetime
3. from airflow import DAG
4. from airflow.providers.docker.operators.docker import DockerOperator
5. from airflow.sensors.filesystem import FileSensor
6.
7. from docker.types import Mount
8.
9. default_args = {
10.     'owner': 'airflow',
11.     'start_date': datetime(2023, 1, 1),
12.     'retries': 1,
13. }
14.
15.
16. dag = DAG(
17.     'audio_to_text_to_summary_to_pdf',
18.     default_args=default_args,
19.     description='DAG for extracting audio, transforming to text, summarizing, and saving
    as PDF',
```

```

20.     schedule_interval=None,
21. )
22.
23. # TODO: Connection could be done via PythonAPI - but I didnt found HOW - so, do this in
    Web instead...
24. #file_connection = Connection(
25. #     conn_id="file_connection",
26. #     conn_type="fs",
27. #     description="Connection to file-path",
28. #)
29.
30. wait_for_new_file = FileSensor(
31.     task_id='wait_for_new_file',
32.     poke_interval=10, # Interval to check for new files (in seconds)
33.     filepath='/opt/airflow/data/lab2', # Target folder to monitor
34.     fs_conn_id='file_connection',
35.     dag=dag,
36. )
37.
38. extract_audio = DockerOperator(
39.     task_id='extract_audio',
40.     image='jrottenberg/ffmpeg',
41.     docker_url="tcp://docker:2375", # For Dind usage case
42.     mount_tmp_dir=False,
43.     network_mode='bridge',
44.     entrypoint='bash',
45.     command=['-c', 'cd /data/lab2 && for single_video in ./*.mp4; do ffmpeg -y -i
    "${single_video}" -ss 1 -to 5 -vn "../lab2_output/${single_video}.wav"; done'],
46.     mounts=[
47.         Mount(source='/data', target='/data', type='bind'),
48.     ],
49.     dag=dag,
50. )
51.
52. audo2text = DockerOperator(
53.     task_id='audio2text',
54.     image='huggy_face_image',
55.     docker_url="tcp://docker:2375", # For Dind usage case
56.     mount_tmp_dir=False,
57.     network_mode='bridge',
58.     entrypoint='bash',
59.     command=['-c', "python /data/audio2text.py"],
60.     mounts=[
61.         Mount(source='/data', target='/data', type='bind'),
62.     ],
63.     dag=dag,
64. )
65.
66. text2summary = DockerOperator(
67.     task_id='text2summary',
68.     image='huggy_face_image',
69.     docker_url="tcp://docker:2375", # For Dind usage case
70.     mount_tmp_dir=False,
71.     network_mode='bridge',
72.     entrypoint='bash',
73.     command=['-c', "python /data/text2summary.py"],
74.     mounts=[
75.         Mount(source='/data', target='/data', type='bind'),
76.     ],
77.     dag=dag,
78. )
79.
80. wait_for_new_file >> extract_audio >> audo2text >> text2summary

```

Kod audio2text:

```

1. from transformers import pipeline

```

```

2. import glob
3. import os
4.
5. for audio_file_path in glob.glob('/data/lab2_output/*.wav'):
6.     _, filename = os.path.split(audio_file_path)
7.     output_txt_file_path = f'/data/lab2_output/text/{filename}.txt'
8.     if os.path.isfile(output_txt_file_path):
9.         continue # skip, already exist
10.
11.     pipe = pipeline("automatic-speech-recognition", "openai/whisper-tiny")
12.     res = pipe(audio_file_path)
13.
14.     with open(output_txt_file_path, "w") as text_file:
15.         text_file.write(res['text'])

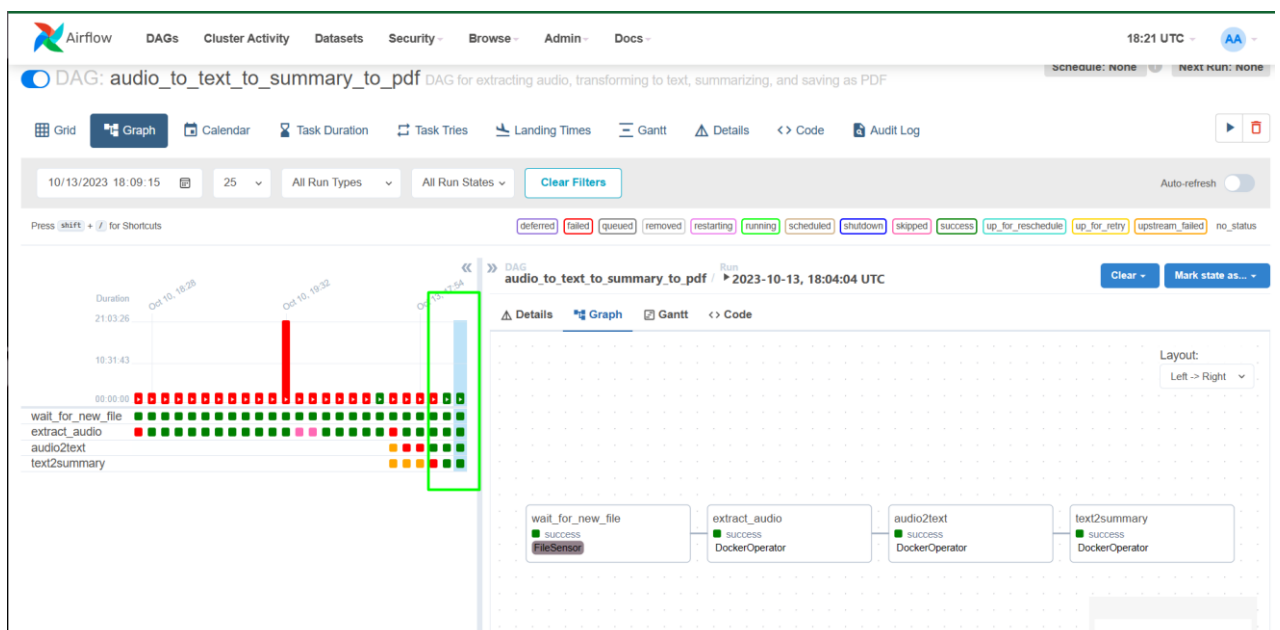
```

Kod text2summary:

```

1. from transformers import pipeline
2. import glob
3. import os
4.
5. for txt_file_path in glob.glob('/data/lab2_output/text/*.txt'):
6.     _, filename = os.path.split(txt_file_path)
7.     output_summary_file_path = f'/data/lab2_output/summary/{filename}'
8.     if os.path.isfile(output_summary_file_path):
9.         continue # skip, already exist
10.
11.     with open(txt_file_path, 'r') as fr:
12.         text = fr.read()
13.
14.     summarizer = pipeline("summarization", max_length=9) # Since our input small
15.     text_summ = summarizer(text)
16.
17.     with open(output_summary_file_path, "w") as text_file:
18.         text_file.write(text_summ[0]['summary_text'])

```



Обучение сети при обнаружении нового файла на примере MNIST

Для тестирования такой ситуации, я разбил обучающую выборку MNIST на 10 случайных частей. В папке `airflow/data/lab2_nn_train/full_data` – полные файлы данных, сама загрузка и слежка идет за папкой `airflow/data/lab2_nn_train/data`, куда можно закидывать или удалять файлы для тестирования.

В качестве обучения взят набор данных MNIST и пару легких сверточных сетей.

`Airflow/dags/airflow_train_lab2.py`

```
1. import os
2. from datetime import datetime
3. from airflow import DAG
4. from airflow.providers.docker.operators.docker import DockerOperator
5. from airflow.sensors.filesystem import FileSensor
6.
7. from docker.types import Mount
8.
9. default_args = {
10.     'owner': 'airflow',
11.     'start_date': datetime(2023, 1, 1),
12.     'retries': 1,
13. }
14.
15.
16. dag = DAG(
17.     'train_nn',
18.     default_args=default_args,
19.     description='DAG train NN',
20.     schedule_interval=None,
21. )
22.
23. wait_for_new_file = FileSensor(
24.     task_id='wait_for_new_train_file',
25.     poke_interval=10, # Interval to check for new files (in seconds)
26.     filepath='/opt/airflow/data/lab2_nn_train/data', # Target folder to monitor
27.     fs_conn_id='file_train_connection',
28.     dag=dag,
29. )
30.
31. train_nn = DockerOperator(
32.     task_id='train_nn_on_updated_data',
33.     image='huggy_face_image',
34.     docker_url="tcp://docker:2375", # For Dind usage case
35.     mount_tmp_dir=False,
36.     network_mode='bridge',
37.     entrypoint='bash',
38.     command=['-c', "python /data/lab2_nn_train/train_nn_mnist.py"],
39.     mounts=[
40.         Mount(source='/data', target='/data', type='bind'),
41.     ],
42.     dag=dag,
43. )
44.
```

```
45.  
46. wait_for_new_file >> train_nn
```

Код обучения:

```
1. import argparse  
2. import torch  
3. import torch.nn as nn  
4. import torch.nn.functional as F  
5. import torch.optim as optim  
6. from torch.utils.data import Dataset  
7. from torchvision import transforms  
8. from torch.optim.lr_scheduler import StepLR  
9.  
10. import numpy as np  
11. import glob  
12. import os  
13. import sys  
14. from datetime import datetime  
15.  
16. SAVE_MODEL_PATH = '/data/lab2_nn_train/models'  
17. DATA_PATH = '/data/lab2_nn_train/data'  
18. TEST_FILENAME = 'test_data.npz'  
19. TRAIN_PREFIX = 'train_'  
20.  
21.  
22. BATCH_SIZE = 32  
23. EPOCHS = 15  
24. LR = 1e-4  
25. LR_STEP = 0.7  
26. CUDA = False  
27. SEED = 2023  
28. LOG_INTERVAL = 10  
29. SAVE_MODEL = True  
30.  
31.  
32. class MnistDataset(Dataset):  
33.  
34.     def __init__(self, x, y, transform=None):  
35.         assert len(x) == len(y)  
36.         self.x = x  
37.         self.y = y  
38.         self.transform = transform  
39.  
40.     def __len__(self):  
41.         return len(self.x)  
42.  
43.     def __getitem__(self, idx):  
44.         x_sample = self.x[idx]  
45.  
46.         if self.transform:  
47.             x_sample = self.transform(x_sample)  
48.  
49.         return x_sample, self.y[idx]  
50.  
51.  
52. class Net(nn.Module):  
53.     def __init__(self):  
54.         super(Net, self).__init__()  
55.         self.conv1 = nn.Conv2d(1, 32, 3, 1)  
56.         self.conv2 = nn.Conv2d(32, 64, 3, 1)  
57.         self.dropout1 = nn.Dropout(0.25)  
58.         self.dropout2 = nn.Dropout(0.5)  
59.         self.fc1 = nn.Linear(9216, 128)  
60.         self.fc2 = nn.Linear(128, 10)  
61.  
62.     def forward(self, x):
```

```

63.         x = self.conv1(x)
64.         x = F.relu(x)
65.         x = self.conv2(x)
66.         x = F.relu(x)
67.         x = F.max_pool2d(x, 2)
68.         x = self.dropout1(x)
69.         x = torch.flatten(x, 1)
70.         x = self.fc1(x)
71.         x = F.relu(x)
72.         x = self.dropout2(x)
73.         x = self.fc2(x)
74.         output = F.log_softmax(x, dim=1)
75.         return output
76.
77.
78. def train(model, device, train_loader, optimizer, epoch):
79.     model.train()
80.     for batch_idx, (data, target) in enumerate(train_loader):
81.         data, target = data.to(device), target.to(device)
82.         optimizer.zero_grad()
83.         output = model(data)
84.         loss = F.nll_loss(output, target)
85.         loss.backward()
86.         optimizer.step()
87.         if batch_idx % LOG_INTERVAL == 0:
88.             print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
89.                 epoch, batch_idx * len(data), len(train_loader.dataset),
90.                 100. * batch_idx / len(train_loader), loss.item()))
91.
92.
93. def test(model, device, test_loader):
94.     model.eval()
95.     test_loss = 0
96.     correct = 0
97.     with torch.no_grad():
98.         for data, target in test_loader:
99.             data, target = data.to(device), target.to(device)
100.            output = model(data)
101.            test_loss += F.nll_loss(output, target, reduction='sum').item() #
sum up batch loss
102.            pred = output.argmax(dim=1, keepdim=True) # get the index of the max
log-probability
103.            correct += pred.eq(target.view_as(pred)).sum().item()
104.
105.        test_loss /= len(test_loader.dataset)
106.
107.        print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)\n'.format(
108.            test_loss, correct, len(test_loader.dataset),
109.            100. * correct / len(test_loader.dataset)))
110.
111.
112. def main():
113.     use_cuda = CUDA and torch.cuda.is_available()
114.
115.     torch.manual_seed(SEED)
116.     exp_folder_path = os.path.join('/data/lab2_nn_train/res',
str(datetime.now()))
117.     os.makedirs(exp_folder_path, exist_ok=True)
118.     sys.stdout = open(os.path.join(exp_folder_path, "print.log"), 'w')
119.
120.     if use_cuda:
121.         device = torch.device("cuda")
122.     else:
123.         device = torch.device("cpu")
124.
125.     train_kwargs = {'batch_size': BATCH_SIZE}
126.     test_kwargs = {'batch_size': BATCH_SIZE}
127.     if use_cuda:
128.         cuda_kwargs = {'num_workers': 1,

```

```

129.                 'pin_memory': True,
130.                 'shuffle': True}
131.         train_kwargs.update(cuda_kwargs)
132.         test_kwargs.update(cuda_kwargs)
133.
134.         transform=transforms.Compose([
135.             transforms.ToTensor(),
136.             transforms.Normalize((0.1307,), (0.3081,))
137.         ])
138.         x_train_list = []
139.         y_train_list = []
140.
141.         for single_train_file in glob.glob(os.path.join(DATA_PATH,
142. f'{TRAIN_PREFIX}*.npz')):
143.             single_train_loaded = np.load(single_train_file)
144.             x_train_list.append(single_train_loaded['x_train'])
145.             y_train_list.append(single_train_loaded['y_train'])
146.
147.         x_train_np = np.concatenate(x_train_list, axis=0)
148.         y_train_np = np.concatenate(y_train_list, axis=0)
149.         dataset_train = MnistDataset(
150.             x=x_train_np, y=y_train_np,
151.             transform=transform
152.         )
153.         test_loaded = np.load(os.path.join(DATA_PATH, TEST_FILENAME))
154.         dataset_test = MnistDataset(
155.             x=test_loaded['x_test'], y=test_loaded['y_test'],
156.             transform=transform
157.         )
158.         train_loader = torch.utils.data.DataLoader(dataset_train,**train_kwargs)
159.         test_loader = torch.utils.data.DataLoader(dataset_test, **test_kwargs)
160.
161.         model = Net().to(device)
162.         optimizer = optim.Adadelta(model.parameters(), lr=LR)
163.
164.         scheduler = StepLR(optimizer, step_size=1, gamma=LR_STEP)
165.         for epoch in range(1, EPOCHS + 1):
166.             train(model, device, train_loader, optimizer, epoch)
167.             test(model, device, test_loader)
168.             scheduler.step()
169.
170.         if SAVE_MODEL:
171.             torch.save(model.state_dict(), os.path.join(exp_folder_path,
172. 'mnist_cnn.pt'))
173.
174.         if __name__ == '__main__':
175.             main()

```


Airflow

DAGs

Cluster Activity

Datasets

Security

Browse

Admin

Docs

18:36 UTC

AA

DAG: train_nn DAG train NN

Schedule: None

Next Run: None

Grid

Graph

Calendar

Task Duration

Task Tries

Landing Times

Gantt

Details

Code

Audit Log

10/13/2023 18:19:21

25

All Run Types

All Run States

Clear Filters

Auto-refresh

Press **shift** + **/** for Shortcuts

deferred

failed

queued

removed

restarting

running

scheduled

shutdown

skipped

success

up_for_reschedule

up_for_retry

upstream_failed

no_status

Duration

00:15.44

00:07.52

00:00.00

wait_for_new_train_file

train_nn_on_updated_data

0d 16:28:12

0d 16:28:12

DAG

Run

Task

train_nn

2023-10-13, 18:19:20 UTC

train_nn_on_updated_data

Clear task

Mark state as...

Filter Tasks

Details

Graph

Gantt

Code

Logs

(by attempts)

1

All Levels

All File Sources

Wrap

Download

See More

8b7b7698de32

*** found local files:

*** + /opt/airflow/logs/dag_id=train_nn/run_id=manual__2023-10-13T18:19:20.549959400:00/task_id=train_nn_on_updated_data/attempt=1.log

[2023-10-13, 18:19:26 UTC] (taskinstance.py:1159) INFO - Dependencies all met for dep_context=non-requeueable deps ti=<TaskInstance: train_nn.train_nn_on_updated_data manual__2023-10-13T18:19:20.549959400>

[2023-10-13, 18:19:26 UTC] (taskinstance.py:1159) INFO - Dependencies all met for dep_context=requeueable deps ti=<TaskInstance: train_nn.train_nn_on_updated_data manual__2023-10-13T18:19:20.549959400>

[2023-10-13, 18:19:26 UTC] (taskinstance.py:1361) INFO - Starting attempt 1 of 2

[2023-10-13, 18:19:26 UTC] (taskinstance.py:1362) INFO - Executing <Task(DockerOperator): train_nn_on_updated_data> on 2023-10-13 18:19:20.549959400:00

[2023-10-13, 18:19:26 UTC] (standard_task_runner.py:57) INFO - Started process 905 to run task

[2023-10-13, 18:19:26 UTC] (standard_task_runner.py:84) INFO - Running: ['***', 'tasks', 'run', 'train_nn', 'train_nn_on_updated_data', 'manual__2023-10-13T18:19:20.549959400:00', '--job-id', '43']

[2023-10-13, 18:19:26 UTC] (standard_task_runner.py:85) INFO - Job 433: Subtask train_nn_on_updated_data

[2023-10-13, 18:19:26 UTC] (task_command.py:415) INFO - Running (TaskInstance: train_nn.train_nn_on_updated_data manual__2023-10-13T18:19:20.549959400:00 [running]) on host 8b7b7698de32

[2023-10-13, 18:19:26 UTC] (taskinstance.py:1666) INFO - Exporting env vars: AIRFLOW_CTX_DAG_OWNER=*** AIRFLOW_CTX_DAG_ID=train_nn AIRFLOW_CTX_TASK_ID=train_nn_on_updated_data AIRFLOW_CTX_E

[2023-10-13, 18:19:26 UTC] (docker.py:333) INFO - Starting docker container from image huggy_face_image

[2023-10-13, 18:33:25 UTC] (taskinstance.py:1400) INFO - Marking task as SUCCESS. dag_id=train_nn, task_id=train_nn_on_updated_data, execution_date=20231013T181920, start_date=20231013T181926, en

[2023-10-13, 18:33:25 UTC] (local_task_job_runner.py:128) INFO - Task exited with return code 0

[2023-10-13, 18:33:25 UTC] (taskinstance.py:2784) INFO - 0 downstream tasks scheduled from follow-on schedule check